

Programmazione Avanzata per il Calcolo
Scientifico
Advanced Programming for Scientific Computing
Lecture title: Move semantic

Luca Formaggia

MOX
Dipartimento di Matematica
Politecnico di Milano

A.A. 2014/2015

Introduction

lvalue and rvalue examples

The need of move semantic

rvalue reference

- Overloading rules

- Forcing a move

The perfect forwarding problem

Introduction

One of the problems of C++11 is that often objects can be of big size (think of a matrix for instance). Thus, we should avoid to make useless copies. Unfortunately, copies may happen in different places. Let's for instance look at this piece of software that swaps two matrices.

```
Matrix a,b;  
... //some work with the matrices  
Matrix temp(a);  
a=b;  
b=temp;
```

This is memory inefficient, we do not really need to have at the same time temp, a and !b!. We just want to swap the contents of the two Matrix objects!

This lecture is largely based on the material found in thbecker.net/articles/rvalue_reference and chapter 5 of the book **Effective modern C++** by H. Sutter.

Move semantic and perfect forwarding

The new standard has taken a decisive step to solve this problems with the introduction of **rvalue references**, which allow to implement:

Move semantics: makes it possible for compilers to replace expensive copying operations with less expensive moves. **Move constructors** and **move assignment operators** offer control over the semantics of moving.

Perfect forwarding makes it possible to write function templates that take arbitrary arguments and forward them to other functions such that the target functions receive exactly the same arguments as were passed to the forwarding functions.

However before dealing with move semantic it is better to revise some important concepts.

types and categories

In C++ expressions and variables are characterized by two different properties: **type** and **category**.

The type identifies an expression in terms of its size and the functionality it provides. A category is instead linked to the context in which the expression or variable is used.

Type

A type may itself be subdivided in basic type, qualifiers and “adornment”.

More than giving definitions, it is simpler to give examples

```
int a; // basic type is int  
const int a; // basic type is int, qualified as const  
int& c=a; // basic type is int. Type is reference to int
```

& is an adornment that specifies that the type is a (lvalue) reference, i.e. an alias to another object.

The two most important qualifiers are **const** and **volatile**. **const** indicates that the variable cannot be changed, **volatile** that it can change between different accesses, even if it does not appear to be modified. **mutable** is another qualifier.

Categories

There are in fact three primary categories in C++11: `lvalue`, `rvalue` and `xvalue`.

But, for the sake of simplicity we identify the last two with their more common name: `rvalue`.

rvalues and lvalues

The original definition of lvalues and rvalues from the earliest days of C is as follows: An lvalue is an expression that may appear on the left or on the right hand side of an assignment, whereas an rvalue is an expression that **can only appear on the right hand side of an assignment.**

```
double fun(); // a function returning a double  
3.14=a; //WRONG a literal expression is a rvalue!  
fun()=5; //WRONG returning an object generates an rvalue
```

Rvalues in C++

User defined types and operator overloading makes the definition of rvalues/lvalues rather complicated in C++. We avoid the formal definition contained in the standard (very technical), which in fact distinguishes rvalues in prvalues and xvalues. We give a simple definition, correct in the 99.8 per cent of cases:

An lvalue is an expression that refers to a memory location and allows us to take the address of that memory location via the & operator. An rvalue is an expression that is not an lvalue.

What does the “l” in lvalue mean?

An rvalue cannot be used to initialize a non-const lvalue reference (while it can be used to initialize const references)

That is, an rvalue cannot be converted to an lvalue, but when an lvalue is used in a context where an rvalue is expected, the lvalue is **implicitly converted to an rvalue**.

```
// lvalues:
int i = 42;
i = 43; // ok, i is an lvalue
int* p = &i; // ok, i is an lvalue
const & g=42; //
int& foo();
foo() = 42; // ok, foo() is an lvalue
int* p1 = &foo(); // ok, foo() is an lvalue
// rvalues:
int foobar();
int j = 0;
j = foobar(); // ok, foobar() is an rvalue
int* p2 = &foobar(); // error
// cannot take the address of an rvalue
j = 42; // ok, 42 is an rvalue
```

Another example

```
struct UserType
{
    double member_function();
}
int var = 0;
var = 1 + 2; // ok, var is an lvalue here
var + 1 = 2 + 3; // error, var + 1 is an rvalue
int* p1 = &var; // ok, var is an lvalue
int* p2 = &(var + 1); // error, var + 1 is an rvalue
const double & a=UserType().member_function(); // ok, calling a
// member function of the class returns a rvalue
// and a rvalue may bind to a const lvalue reference
// THE LIFESPAN OF THE RETURNED OBJECT IS EXTENDED!
```

Reference binding

The process by which an expression is associated to a reference is called binding. The rules for ordinary (lvalue) references are rather straightforward. The interplay between parameter types and binding is used for function overloading.

```
void foo(int & a);  
void foo(const int & a);  
void goo(const int & a);  
...  
foo(5); //calls foo(const int &)  
int g;  
foo (g); // calls foo(const int &)  
goo (g); // goo(const int &);
```

In the call of foo the compiler chooses the best match. In the first case an rvalue const reference (a literal) is bound to a constant reference, in the second a (non const) **int** is bound to a **const int**&. In the call of goo a **int** is bound to a **const int**& since there is no better match.

The need of moving objects

Let's consider a simple example. A function that swaps the arguments

```
template<class T>
void swap (T & a, T& b){
    T tmp{a}; // make a copy of a
    a = b; // copy-assign b to
    b = tmp; // copy assign tmp to b
} // tmp is destroyed on exit.
```

If *a* and *b* are objects of big size this function is memory-inefficient. We have to store tmp.

The optimal swap

We would like to have instead an algorithm of this sort:

- ▶ Move `a` into `tmp`, leaving `a` empty;
- ▶ Fill `a` by moving `b`, leaving `b` empty;
- ▶ Fill `b` by moving `tmp`, leaving `tmp` empty.

It is now possible thanks to the **move semantic**.

rvalue references

To implement move semantic C++11 has introduced a new type, or more precisely a new type decorator, called **rvalue reference**, indicated by `&&`.

Its main usage is in operator overloading: it allows to select operations that move objects instead of copying them.

This is possible because a non-constant (and thus movable) rvalue **preferably binds to a rvalue reference**. But they can bind also to constant lvalue references (so the old behavior is kept if move semantic is not implemented).

While, lvalues can only bind to lvalue (ordinary) references `&`.

Categories of function return values

We also have the following important rules.

- ▶ If a function returns a value that value is considered an **rvalue**.
- ▶ If a function returns a lvalue reference (const or non-const) that value is considered an lvalue.
- ▶ If a function returns a rvalue reference (but there is normally no reason to do so!), that value is an rvalue.

This is fundamental for move semantic.

To understand things better let's look at the simple example in [Bindings/main.cpp](#).

```
void foo(int & a);  
void foo(const int & a);  
int createFive();  
void goo(int & a);  
void goo(int && a);  
void goo(const int & a);
```

I have two overloaded functions and a function returning an **int**.

The main

```
foo(25);           //foo(const int &)  
foo(a);           //foo(int &)  
foo(b);           //foo(int &)  
foo(createFive()); //foo(const int &)  
goo(25);          //goo(int &&) NOTE!  
goo(a);           //goo(int &)  
goo(createFive()); //goo(int &&) NOTE!  
goo(b);           //goo(int &)
```

The general overloading rules

If you implement **void** foo(X&); but not **void** foo(X&&); the behavior is the usual C++98 one: foo can be called on lvalues, but not on rvalues.

If you implement **void** foo(X **const** &); but not **void** foo(X&&); then again, the behavior is the old one: foo can be called on lvalues and rvalues, but it is not possible to distinguish between them.

That is possible only by implementing **void** foo(X&&); as well.

Finally, if you implement **void** foo(X&&); but neither one of **void** foo(X&); and **void** foo(X **const** &); then foo can be called on rvalues, but trying to call it on an lvalue will trigger a compile error.

A first important note

Never ever write a function that returns a reference to a temporary.

```
Matrix & createMatrix(){  
    Matrix tmp;  
    ...  
    return tmp;  
}
```

This not only wrong, but useless! **Temporaries must be always returned by values**. Thanks to RVO (return value optimization) and the fact that the returned object is a rvalue returning by value is optimal!

```
Matrix  createMatrix(); // OK!
```

Another important thing to remember

Named variables of rvalue reference type are lvalues! In particular function parameters (of any function, also constructors) are lvalues!, even if their type is an rvalue reference.

Inside the scope of this function

```
void f(Matrix&& m){  
    ....
```

m is an lvalue.

Remember that the terms *rvalue* and *lvalue* refer to categories, not types. You can take the address of m (i.e. you can write `Matrix * pm=&m`) so it is an lvalue.

How to implement move semantic?

You need to write a move constructor and a move assignment operator.

Beware that unless you have defined (even if with the keyword **default**) come other constructor or copy assignment (see next slide) the compiler provides a move constructor and move assignment operator automatically.

This is the standard signature of move operations for a class named Foo:

```
Foo(Foo&&); // move constructor  
Foo & operator=(Foo&&); // move assignment
```


The rules for the generation of automatic operators

compiler implicitly declares							
user declares		default constructor	destructor	copy constructor	copy assignment	move constructor	move assignment
	Nothing	defaulted	defaulted	defaulted	defaulted	defaulted	defaulted
	Any constructor	not declared	defaulted	defaulted	defaulted	defaulted	defaulted
	default constructor	user declared	defaulted	defaulted	defaulted	defaulted	defaulted
	destructor	defaulted	user declared	defaulted	defaulted	not declared	not declared
	copy constructor	not declared	defaulted	user declared	defaulted	not declared	not declared
	copy assignment	defaulted	defaulted	defaulted	user declared	not declared	not declared
	move constructor	not declared	defaulted	deleted	deleted	user declared	not declared
	move assignment	defaulted	defaulted	deleted	deleted	not declared	user declared

Let us consider the version of MyMat0 class that holds the matrix data in the member **double** * data (not the best choice but good for this example).

A possible constructor and copy-assignment will take the form

```
MyMat0&(MyMat0 const & rhs):data(new double[rhs.nr*rhs.nc])
{
    // make a deep copy
    for (i=0;i<rhs.nr*rhs.nc;++i) data[i]=rhs.data[i];
}
MyMat0& operator=(MyMat0 const & rhs){
    // release the resource
    delete [] this→data;
    // make a deep copy
    data=new double[rhs.nr*rhs.nc];
    for (i=0;i<rhs.nr*rhs.nc;++i) data[i]=rhs.data[i];
}
```

The move operators

The corresponding move operator are

```
MyMat0(MyMat0&& rhs):nr(rhs.nr),nc(rhs.nc)
{
    //get the resource
    this->data=rhs.data;
    rhs.data=nullptr;
    rhs.nc=rhs.nr=0;
}
MyMat0& operator=(MyMat0 const & rhs){
    delete[] this->data; // release the resource
    // make a deep copy
    data=new double[rhs.nr*rhs.nc];
    for (i=0;i<rhs.nr*rhs.nc;++i) data[i]=rhs.data[i];
    rhs.data=nullptr;
    rhs.nc=rhs.nr=0;
}
```

I just grab the resource!

it is important to ensure that the moved object can be deleted correctly!. Since the destructor of MyMat0 calls **delete**[] on data, I

The consequence

```
MyMat0 foo();
```

```
...
```

```
MyMat0 a(foo()); // move constructor is called!
```

```
a=foo(); // move assignement is called
```

Forcing a move

If named variables, and consequently function parameters, are lvalues, how can I apply the move semantic on them? Well, two utilities of the standard library come to our rescue: `std::move()` and `std::forward<T>()`

`std::move(expr)` unconditionally casts `expr` to a rvalue reference
`std::forward<T>(expr)` casts `expr` to the lvalue reference `T&` if `expr` is an lvalue, and to the rvalue reference `T&&` if `expr` is an rvalue.

We will see the latter in the context of perfect forwarding and universal references (another feature of C++11). For the moment let's concentrate on `move()`.

A new version of swap

```
template<class T>
void swap(T& a, T& b) {
    T tmp(std::move(a));
    a = std::move(b);
    b = std::move(tmp);}

```

It type T implements move semantic the swap can be made with much less memory requirement!

Some details

Suppose that a Base class has implemented

```
Base(Base const & rhs); // copy constr.  
Base(Base&& rhs); // move constr.
```

and we want to use the move-constructor of Base in the move constructor of a Derived class. **We need to use std::move**

```
Derived(Derived&& rhs) : Base(std::move(rhs)){  
    // Derived-specific stuff  
}
```

This because rhs is an lvalue (it has a name!). Without move(), Base(rhs) would use the copy-constructor of Base !!



I like to move it. Move it!

All std containers support move semantic (since C++11) and all std algorithms have been rewritten so that if the contained type implements move semantic the creation of unnecessary temporaries can be avoided.

For instance, `std::swap()` does exactly what we have seen in a previous slide and `std::sort()` (which does a lot of swaps) will now be much more efficient on big objects.

Move semantic will also make a few (not all!) template metaprogramming techniques now used in some libraries (like the Eigen) to avoid large size temporaries unnecessary (with a considerable reduction of headaches).

An example

In [MoveSemantic](#) an example of a class implementing a full matrix where we have defined move constructor and move assignment operator.

After compiling it do `make test` and see the result of the memory usage of the version of the code where move semantic is activated and the one where is deactivated (by defining the preprocessor macro `NOMOVE`).

A perfect forwarding

There is another problem that rvalue references can solve. Let's look at this function that implements the object factory design pattern

```
template<typename Arg>
unique_ptr<Base> factory(Arg arg, int switch){
    if(switch==1)
        return unique_ptr<Base>(new D1(arg));
    ... etc}
```

We are making a useless copy of the first argument into the parameter arg, while what we want is just to forward arg to the constructor of D1

One may think that the problem is solved by using references. But this solution has some other drawbacks, we will not detail them, see [here](#) for a discussion. Furthermore all “standard” solutions will block move semantic.

We do not want to go into too much detail, we just present the C++11 solution. But first a digression on [universal references](#).

Universal references

First of all there is no mention of universal references in the C++11 standard. It is a term invented by H. Sutter to explain in simple terms a particular behavior of rvalue references when used as template dependent parameters.

This behavior is linked to the so called [reference collapsing rule](#), but if you understand universal references you can avoid the headache of understanding reference collapsing. Universal references appear when you have constructs of the type

```
template <class T>  
double fun(T&& x);
```

How does it work?

The combination of template parameter deduction and collapsing rule causes parameter `x` to be able to **bind both to lvalues and rvalues!**

```
double randomValue(); // a function  
double a{5.0};  
double const & ra=a;  
...  
x=fun(6.7); // fun(double&&) (T=double)  
x=fun(a); // fun(double&) (T=double&)  
x=fun(randomValue()); // fun(double&&) (T=double)  
x=fun(ra); // fun(const double&) (T=const double&)
```

That's why they are called universal!.

BEWARE

A rvalue reference behaves as universal reference only if its type is deduced at the moment of the instance of the function!

```
template<class P>  
void fun(P && x); // Universal reference  
template <class T>  
void foo(std::vector<T>&& x); // NOT UNIVERSAL
```

In the second function the rvalue reference is not a universal reference (i.e. it just binds to rvalues) since at the moment of the instance the type is known (well you have to think a little on how templates are instantiated to understand why...)

In other words, universal references takes the (almost) just the form

```
template  
<class T> f(T&& x)
```

(of course the template parameter may have a different name and you may have more than one parameter...).

Perfect forwarding

Now we can use the `std::forward<T>()` function to solve the problem:

```
template<typename Arg>
unique_ptr<Base> factory(Arg&& arg, int switch){
    if(switch==1)
        return unique_ptr<Base>(
            new D1(std::forward<Arg>(arg)));
    ... etc}
```

Thanks to the “universal binding” of universal references and the magic of `std::forward<T>()`, this version works both if `Arg` is an lvalue or rvalue and it resolves rvalue references correctly: no useless temporaries if move semantic is implemented, all automagically.

H. Sutter docet

- ▶ Use `std::move()` if you want to move lvalues, but
- ▶ always use `std::forward<T>` on universal references!

But remember that after a move the moved object is “empty”!

Some useful usage of universal refs

```
struct Foo{  
    template <class T>  
    Foo(T && x):M_x(std::forward<T>(x)){}  
    private:  
    Matrix M_x;  
}
```

A constructor that takes **any argument convertible to a Matrix**. In one go we have both the copy and move version!!!

```
Matrix HilbertMatrix(); // a function  
Matrix m;  
Foo g(m); // copy  
Foo z(HilbertMatrix()); // moved!
```


Cheating the compiler

```
template<class T>
struct Pippo{
    void f(T &&); // NO UNIVERSAL
}
```

It is not universal because when an object of type Pippo<T> is instantiated T assumes a value, so there is nothing to be deduced when calling Pippo<T>::f(). Workaround:

```
template<class T>
struct Pippo{
    template <class Q>
    void f(Q &&); // UNIVERSAL!!
}
```

Conclusions

Move semantic is one of the most relevant features of C++11. It is not simple (but not overly complicated after all).

Do implement move semantic on class of large size. You just have to add the move constructor and the move assignment operator. Remember that the compiler may give you one automatically (but remember also that sometimes it does not). If the automatic move constructor is what you need and you want to be sure to have it, use the `default` keyword.